

PM566 Lab 1

AUTHOR

Chih-Chan Jessica Lan

Input Libray and Data

```
library(datasauRus)
library(ggplot2)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(tidyr)

df <- datasaurus_dozen
```

Question 1

Based on the help file, how many rows and how many columns does the `datasaurus_dozen` file have? What are the variables included in the data frame? Add your responses to your lab report, with relevant code in the associated R code chunk, and free-form text outside of the code chunk.

Answer:

There are 1846 rows and 3 columns in the `datasaurus_dozen` file. The 3 variable names include: `dataset`, `x`, and `y`

```
dim(df)
```

```
[1] 1846    3
```

```
# cat("Dimension of dataset:", dim(df), "\n")
# print('Names of columns in the dataset:', names(df))
```

Summary of dataset names: There are 13 datasets

```
table(datasaurus_dozen$dataset)
```

away	bullseye	circle	dino	dots	h_lines	high_lines
142	142	142	142	142	142	142
slant_down	slant_up	star	v_lines	wide_lines	x_shape	
142	142	142	142	142	142	

Summary of Continuous Variables (x and y)

```
# Variable x
summary(df$x)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
15.56	41.07	52.59	54.27	67.28	98.29

```
# Variable y
summary(df$y)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.01512	22.56107	47.59445	47.83510	71.81078	99.69468


Question 2

Plot y vs. x for the dino dataset. Then, calculate the correlation coefficient between x and y for just this dataset.


Plot

```
dino_data <- df[df$dataset == 'dino', ]
```

```
plot(dino_data$x, dino_data$y)
```




```
ggplot(data = dino_data, mapping = aes(x = x, y = y)) +  
  geom_point() +  
  theme_bw()
```



Reference: <https://stackoverflow.com/questions/10861773/remove-grid-background-color-and-top-and-right-borders-from-ggplot2>

```
dino_data_eye <- dino_data %>%
  mutate(
    eye_category = case_when(
      x >= 40 & x <= 60 & y >= 75 & y <= 88 ~ "Eye",
      TRUE ~ "Other"
    )
  )

ggplot(dino_data_eye, aes(x = x, y = y, color = eye_category)) +
  geom_point() +
  scale_color_manual(values = c("Eye" = "red", "Other" = "gray70")) +
  theme_bw()
```



Calculate correlation coefficient (r)

```
cor(dino_data$x, dino_data$y)
```

```
[1] -0.06447185
```

```
# dino_data |>  
#   summarize(r = cor(x, y))
```


Question 3

Now try it on your own! Plot y vs. x for the star dataset, another one of the datasaurus_dozen. You can (and should) re-use code we introduced above, just replace the dataset name with the desired dataset. Then, calculate the correlation coefficient between x and y for this dataset. How does this value compare to the r of dino?

```
star_data <- df[df$dataset == "star", ]
```

```
###Plot
```

```
plot(star_data$x, star_data$y)
```



Calculate correlation coefficient (r)

```
cor(star_data$x, star_data$y)
```

```
[1] -0.0629611
```

The correlation coefficients (r) of star and dino data are almost the same around -0.06. Although the rs are similar, the visualization plots actually show a significant difference. This is showing that visualizing data is also critical when we are diving into given data.


Question 4

Plot y vs. x for the circle dataset. You can (and should) reuse code we introduced above, just replace the dataset name with the desired dataset. Then, calculate the correlation coefficient between x and y for this dataset. How does this value compare to the r of dino?

```
circle_data <- df[df$dataset == "circle", ]
```

Plot

```
plot(circle_data$x, circle_data$y)
```



Calculate correlation coefficient (r)

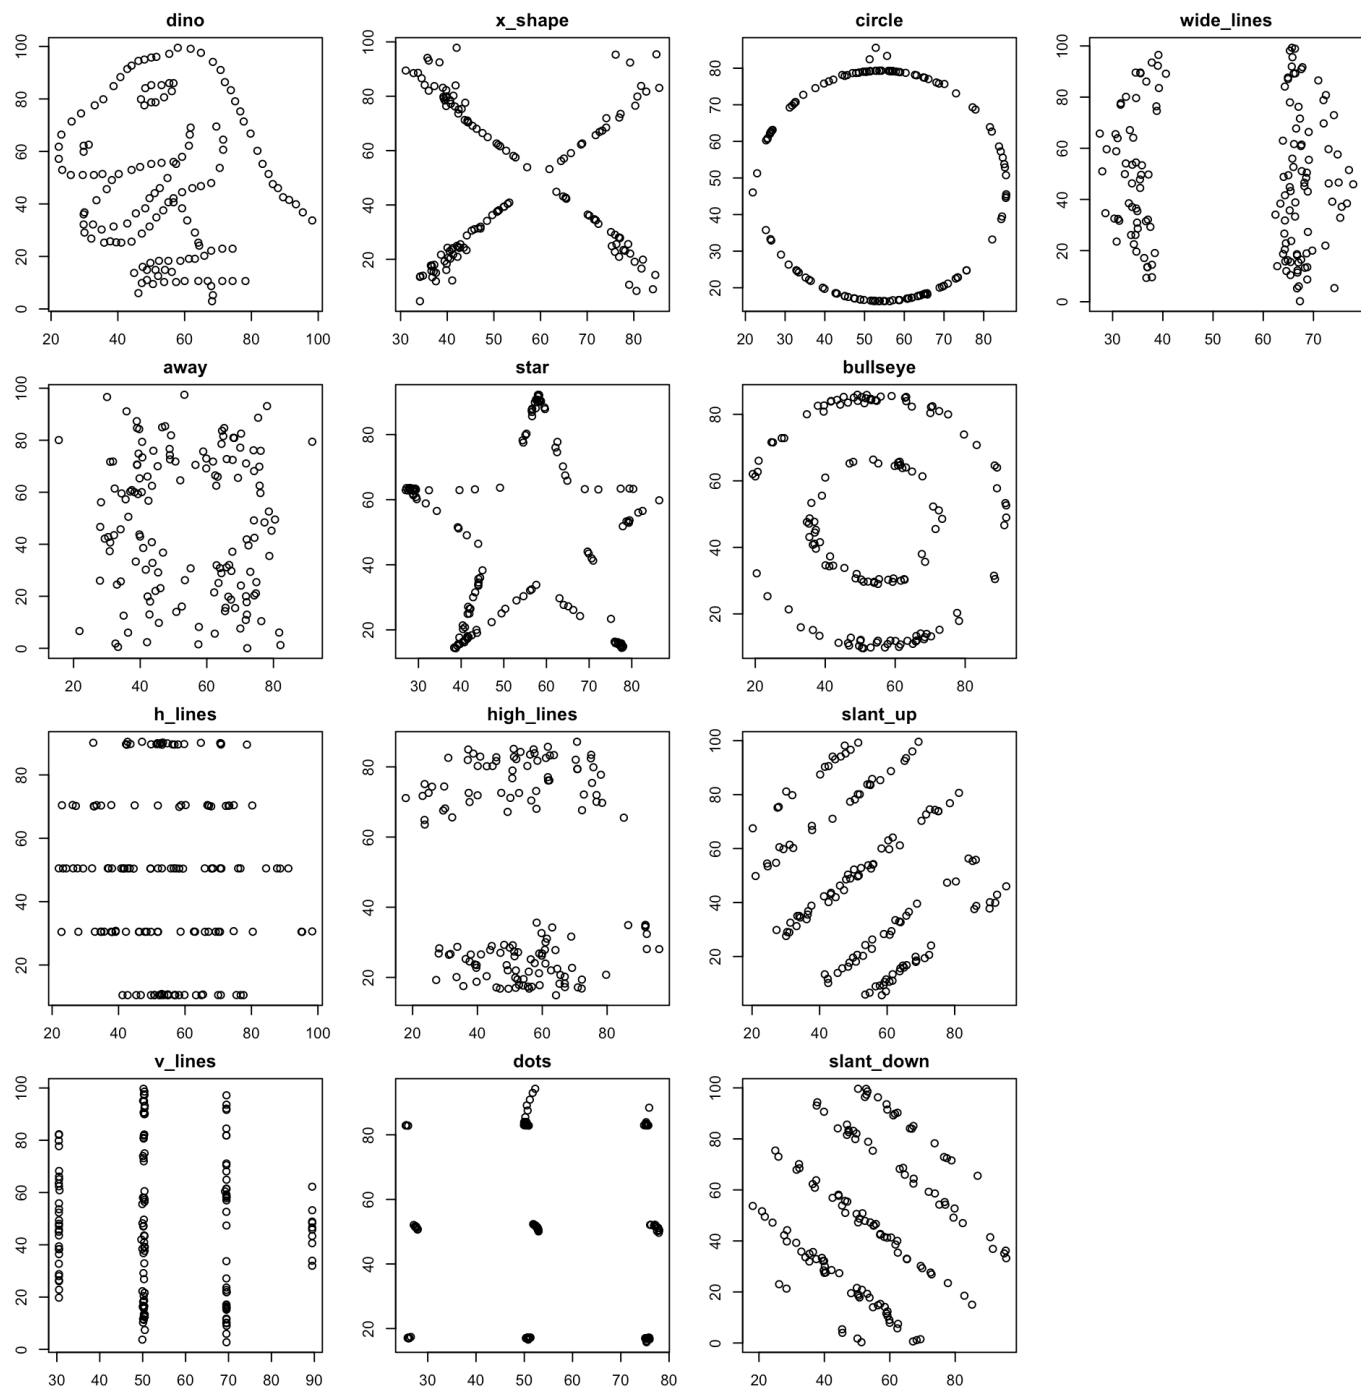
```
cor(circle_data$x, circle_data$y)
```

```
[1] -0.06834336
```

The correlation coefficients (r) of circle and dino data are almost the same around -0.06. Although the rs are similar, the visualization plots actually show a significant difference. This is showing that visualizing data is also critical when we are diving into given data.

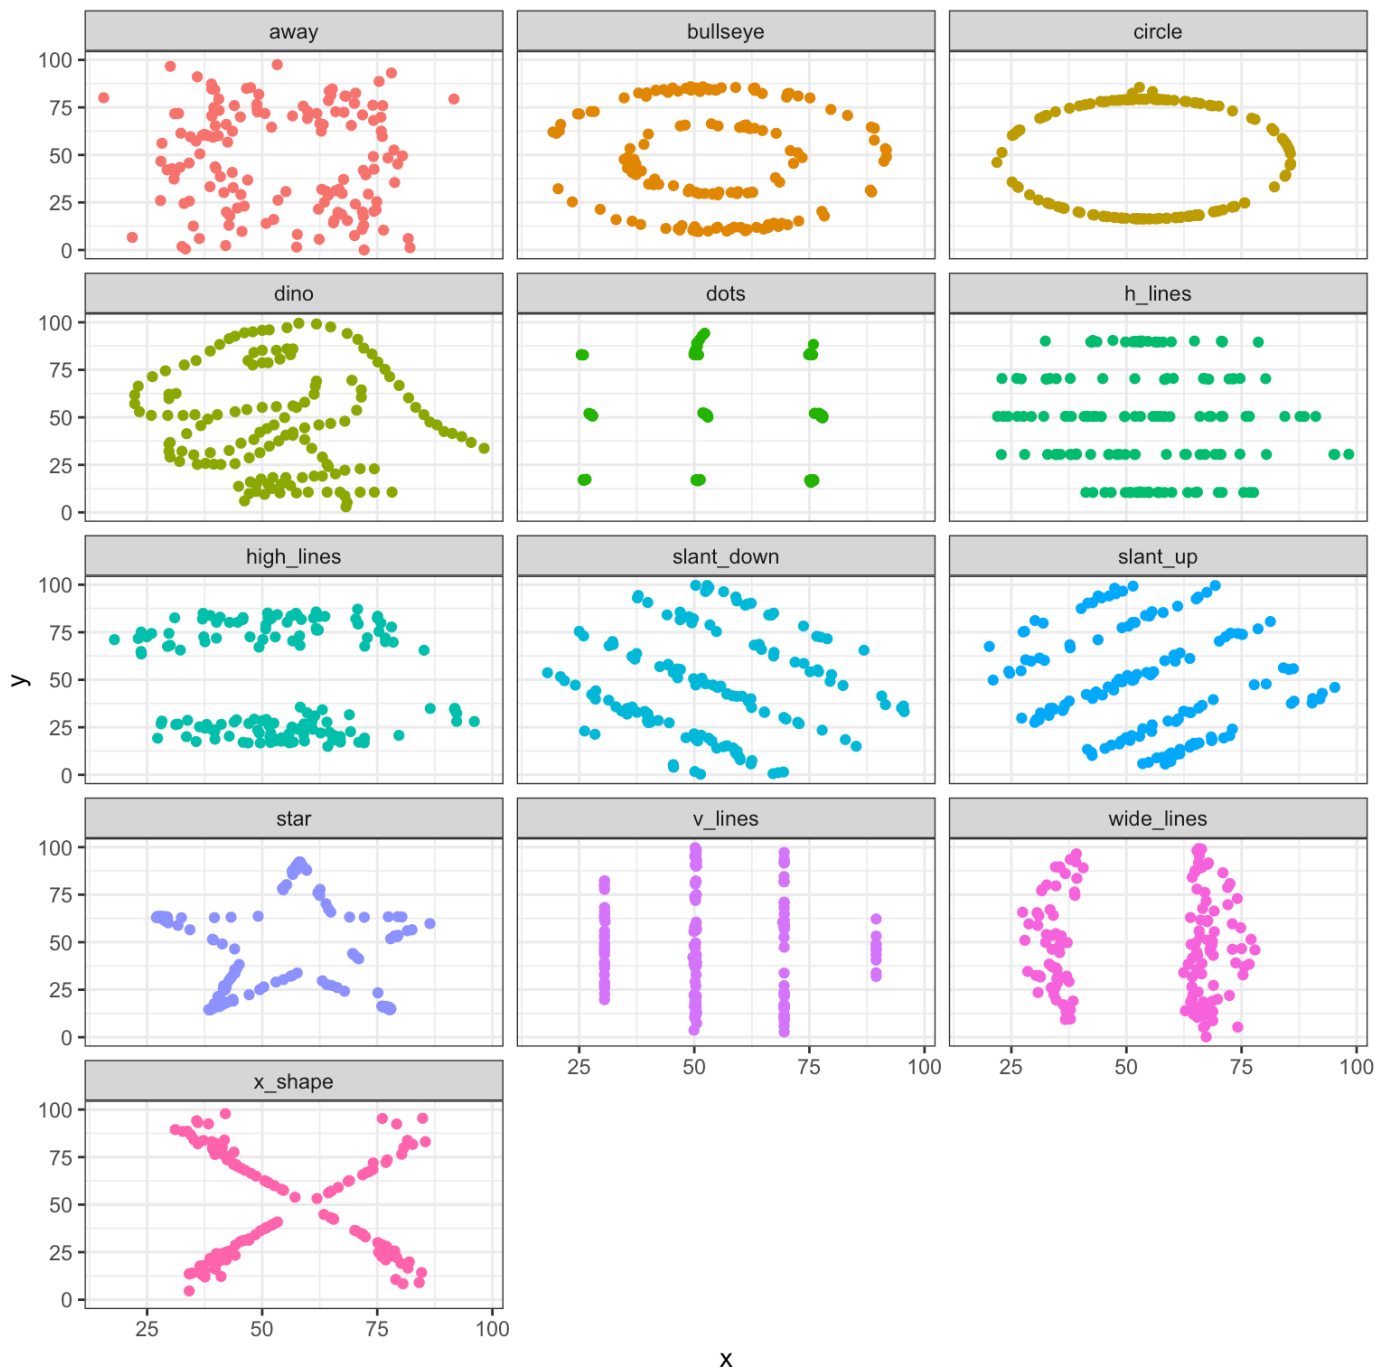
Question 5

```
par(mar = c(2,2,2,2))
layout(matrix(1:16, nrow=4, ncol=4))
for(name in unique(datasaurus_dozen$dataset)){
  subset <- datasaurus_dozen[datasaurus_dozen$dataset == name, ]
  plot(subset$x, subset$y, main = name)
}
layout(1)
```

```
# Assign back to default value
par(mar = c(5,4,4,2) + 0.1)
```

```
ggplot(datasaurus_dozen, aes(x = x, y = y, color = dataset))+
  geom_point() +
  theme_bw() +
  facet_wrap(~ dataset, ncol = 3) +
  theme(legend.position = "none")
```



Question 6

Finally, we want to calculate the correlation between the x and y variables for all 13 datasets. Like before, we will use a loop, but this time, since we want to return a specific value every time through the loop, we will use the `sapply` function. `sapply` is a useful way to apply a function to every element of a vector. In this case, we provide the vector of unique dataset names (like before) and then our own custom function. This function subsets the data as before, and then returns the correlation coefficient as the output of the function.

```
sapply(unique(df$dataset), function(name){  
  subset <- df[df$dataset == name, ]  
  return(cor(subset$x, subset$y))  
})
```

	dino	away	h_lines	v_lines	x_shape	star
	-0.06447185	-0.06412835	-0.06171484	-0.06944557	-0.06558334	-0.06296110
high_lines		dots	circle	bullseye	slant_up	slant_down
	-0.06850422	-0.06034144	-0.06834336	-0.06858639	-0.06860921	-0.06897974
wide_lines						
	-0.06657523					